




The University of New Mexico

---

# Programmable Shaders

Based on the Slides from  
Ed Angel

Professor Emeritus of Computer Science  
University of New Mexico



# Objectives

The University of New Mexico

- 
- Introduce programmable pipelines
    - Vertex shaders
    - Fragment shaders
  - Shader programming with GLSL
  - Related materials
    - Angel: Chapter 2



The University of New Mexico

---

# Programming with OpenGL

## Part 1: Basic Concepts

Modified from the Slides from  
Ed Angel

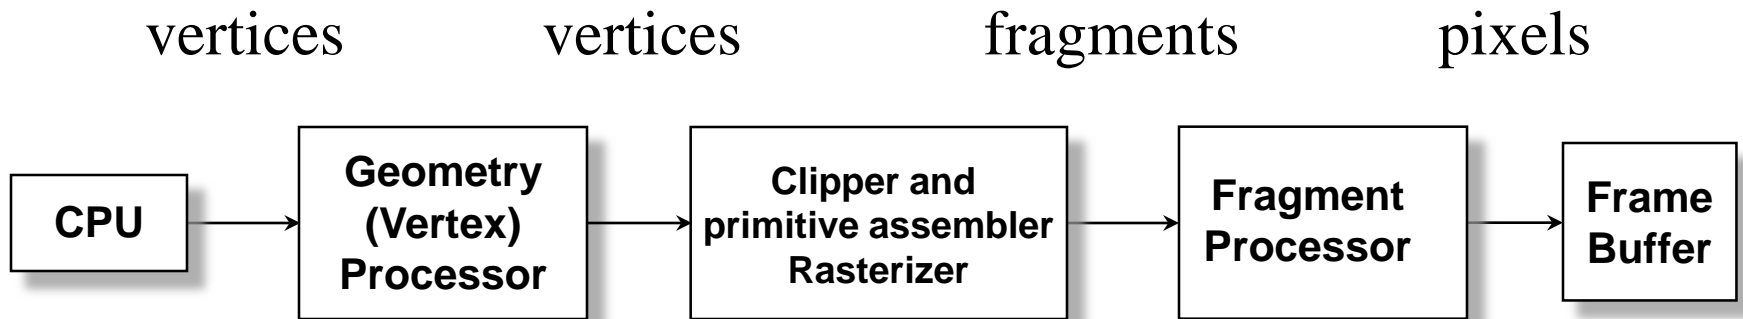
Professor of Emeritus of Computer Science  
University of New Mexico



# Graphics Pipeline

The University of New Mexico

- OpenGL architecture
  - OpenGL as a state machine
  - OpenGL as a data flow machine





# Geometric Calculations

The University of New Mexico

- 
- Geometric data: set of vertices + type
    - Can come from program
    - Type: point, line, polygon
    - Vertex data can be
      - (x,y,z,w) coordinates of a vertex (glVertex)
      - Normal vector
      - Texture Coordinates
      - RGBA color
      - Other data: color indices, edge flags
      - Additional user-defined data in GLSL



# Per-Vertex Operations

The University of New Mexico

- Vertex locations are transformed by the model-view matrix into eye coordinates
- Normals must be transformed with the inverse transpose of the model-view matrix so that  $v \cdot n = v' \cdot n'$  in both spaces
- Textures coordinates are generated if autotexture enabled and the texture matrix is applied



# Primitive Assembly

The University of New Mexico

- 
- Vertices are next assembled into objects
    - Polygons
    - Line Segements
    - Points
  - Transformation by projection matrix
  - Clipping
    - Against user defined planes
    - View volume,  $x=\pm w$ ,  $y=\pm w$ ,  $z=\pm w$
    - Polygon clipping can create new vertices
  - Perspective Division
  - Viewport mapping



# Rasterization

The University of New Mexico

- 
- Geometric entities are rasterized into **fragments**
  - Each fragment corresponds to a point on an integer grid: a displayed pixel
  - Hence each fragment is a *potential pixel*
  - Each fragment has
    - A color
    - Possibly a depth value
    - Texture coordinates






# Fragment Operations

---

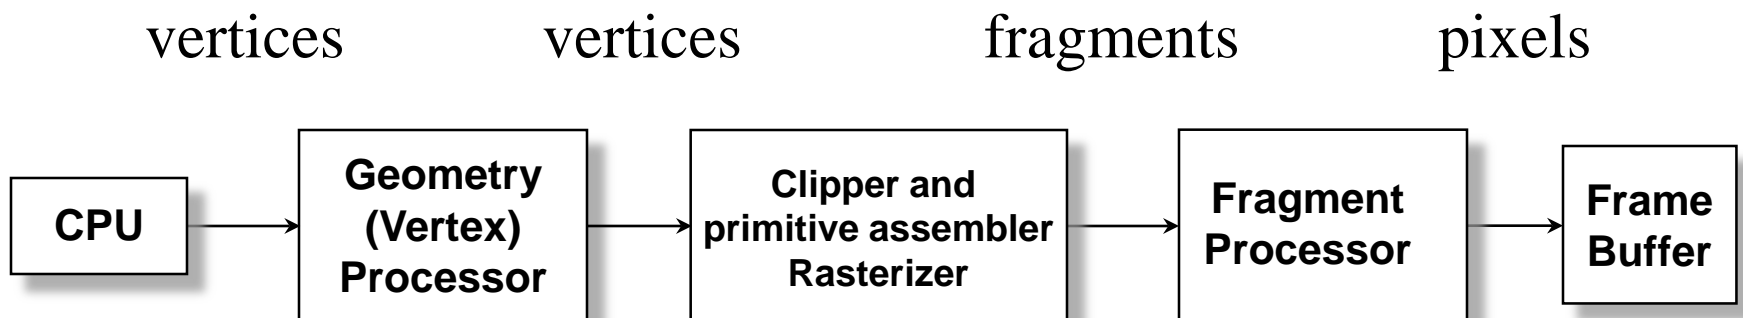
- Texture generation
- Fog
- Antialiasing
- Scissoring
- Alpha test
- Blending
- Dithering
- Logical Operation
- Masking



# OpenGL Architecture

The University of New Mexico

- In obsolete OpenGL, every step in the pipeline is handled by the system (library)
  - OpenGL API
  - Hardware accelerated
  - Unchangeable by the application programmer





# Major Changes

The University of New Mexico

- 
- Recent major advance in real time graphics is a programmable pipeline
    - First introduced by NVIDIA GeForce 3
    - Supported by high-end commodity cards
      - NVIDIA, ATI, 3D Labs
    - Software Support
      - Direct X 8 , 9, 10, 11
      - OpenGL Extensions
      - OpenGL Shading Language (GLSL)
      - High-Level Shading Language (HLSL)
      - Cg



# Programmable Shaders

---

- Replace fixed function vertex and fragment processing by programmable processors called **shaders**
- If we use a programmable shader, we must do all required functions of the fixed function processor



# Various Shaders

The University of New Mexico


- 
- Vertex shader
  - Fragment shader
  - Geometry shader
  - Hull shader, Tessellator, Domain shader
  - Compute shader



# Performance Gains

The University of New Mexico

- Performance is achieved by using GPU rather than CPU
- Control GPU through shader programs
- Application's job is to send data to GPU
- GPU does all rendering and many of geometry/data processing





# OpenGL 3.1

The University of New Mexico

- 
- Totally shader-based
    - No default shaders
    - Each application must provide both a vertex and a fragment shader
  - No immediate mode
  - Few state variables
  - Most 2.5 functions deprecated
  - Backward compatibility not required



# Other Versions

The University of New Mexico

- 
- OpenGL ES
    - Embedded systems
    - Version 1.0 simplified OpenGL 2.1
    - Version 2.0 simplified OpenGL 3.1
      - Shader based
  - WebGL
    - Javascript implementation of ES 2.0
    - Supported on newer browsers
  - OpenGL 4.1 and 4.2
    - Add geometry shaders and tessellator





# What about Direct X?

---

- Windows only
- Advantages
  - Better control of resources
  - Access to high level functionality
- Disadvantages
  - New versions not backward compatible
  - Windows only
- Recent advances in shaders are leading to convergence with OpenGL



# OpenGL and GLSL

The University of New Mexico

- 
- Shader based OpenGL is based less on a state machine model than a data flow model
  - Most state variables, attributes and related pre 3.1 OpenGL functions have been deprecated
  - Action happens in shaders
  - Job of application is to get data to GPU



- 
- OpenGL Shading Language
  - Part of OpenGL 2.0 and up
  - High level C-like language
  - New data types
    - Matrices
    - Vectors
    - Samplers
  - As of OpenGL 3.1, application must provide shaders



The University of New Mexico

---

# Programming with OpenGL

## Part 2: Shaders

Ed Angel

Professor of Emeritus of Computer Science  
University of New Mexico



# Writing Shaders

The University of New Mexico

- 
- First programmable shaders were programmed in an assembly-like manner
    - OpenGL extensions added for vertex and fragment shaders
  - Cg (C for graphics) C-like language for programming shaders
    - Works with both OpenGL and DirectX
    - Complex interface to OpenGL
  - OpenGL Shading Language (GLSL)
  - High-Level Shading Language (HLSL)



- 
- OpenGL Shading Language
  - C-like with
    - Matrix and vector types (2, 3, 4 dimensional)
    - Overloaded operators
    - C++ like constructors
  - Similar to Nvidia's Cg and Microsoft HLSL
  - Code sent to shaders as source code
  - New OpenGL functions to compile, link and get information to shaders



# Vertex Processor

The University of New Mexico

- 
- Takes in vertex data
    - Positions
  - Produces
    - Position in clip coordinates



# Simple Vertex Shader

---

```
in vec4 vPosition;  
void main(void)  
{  
    gl_Position = vPosition;  
}
```

input from application

must link to variable in application

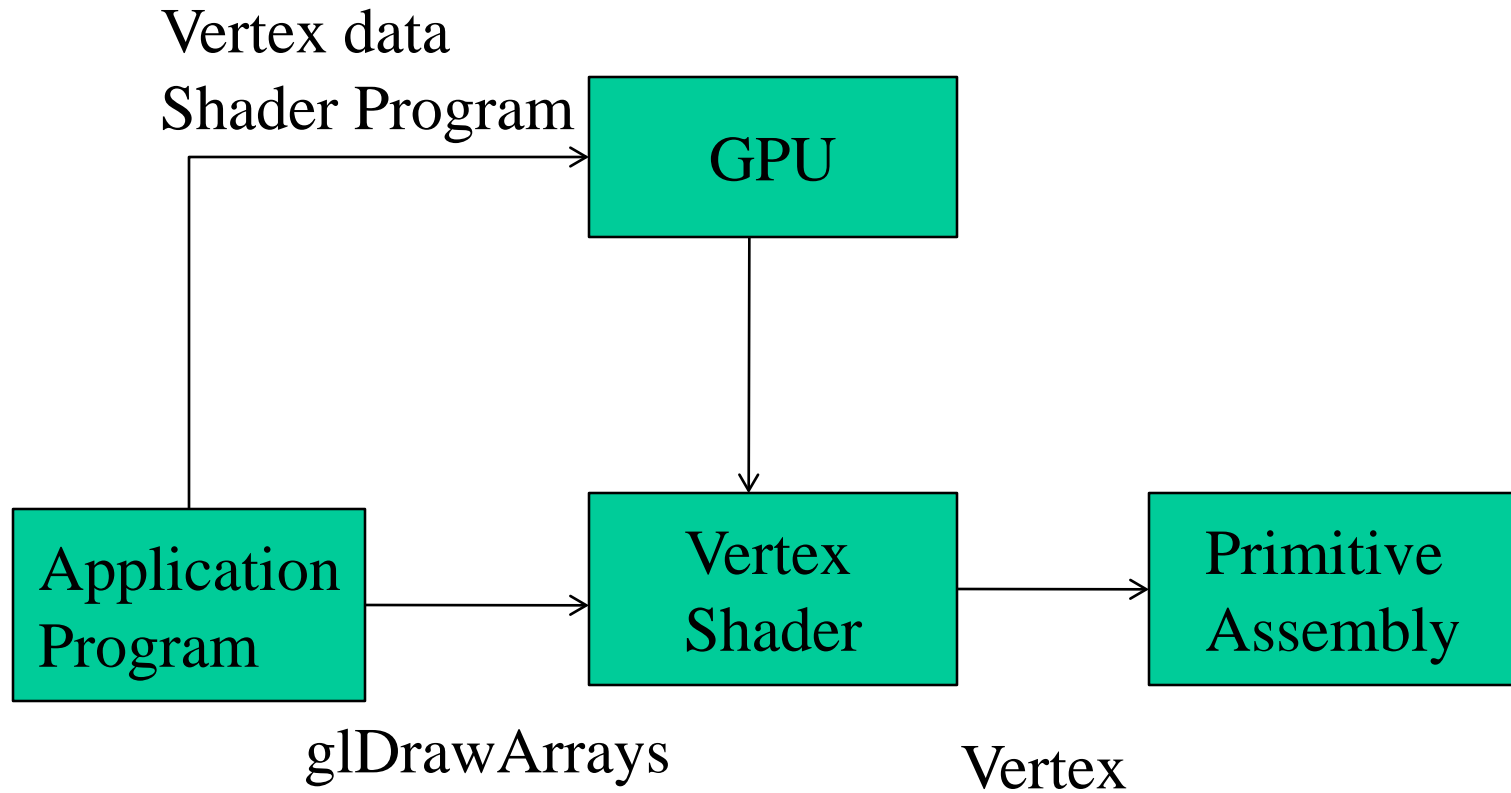
built-in variable






The University of New Mexico

# Execution Model





# Fragment Processor

The University of New Mexico

- 
- Takes in output of rasterizer (fragments)
    - Vertex values have been interpolated over primitive by rasterizer
  - Outputs a fragment
    - Color
  - Fragments still go through fragment tests
    - Hidden-surface removal
    - alpha



The University of New Mexico

# Simple Fragment Shader

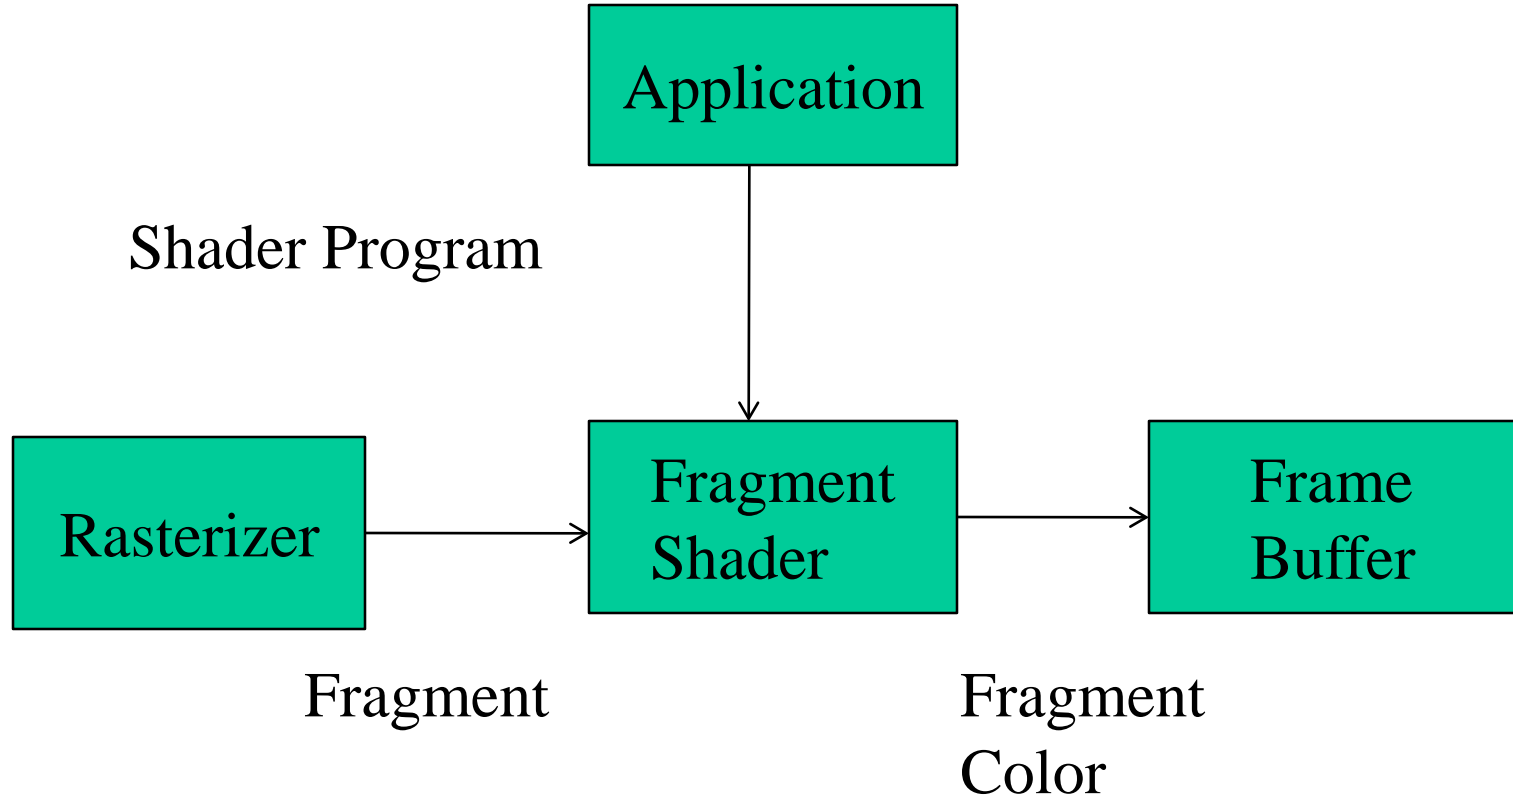
---

```
void main(void)
{
    gl_FragColor = vec4(1.0, 0.0, 0.0, 1.0);
}
```



# Execution Model

---





# Data Types

The University of New Mexico

- 
- C types: int, float, bool
  - Vectors:
    - float vec2, vec3, vec4
    - Also int (ivec) and boolean (bvec)
  - Matrices: mat2, mat3, mat4
    - Stored by columns
    - Standard referencing m[row][column]
  - C++ style constructors
    - `vec3 a = vec3(1.0, 2.0, 3.0)`
    - `vec2 b = vec2(a)`



# Pointers

The University of New Mexico

- 
- There are no pointers in GLSL
  - We can use C structs which can be copied back from functions
  - Because matrices and vectors are basic types, they can be passed into and output from GLSL functions, e.g.

`mat3 func(mat3 a)`



# Qualifiers

The University of New Mexico

- 
- GLSL has many of the same qualifiers such as **const** as C/C++
  - Need others due to the nature of the execution model
  - Variables can change
    - Per primitive (**uniform**)
    - Per vertex (**in**)
    - Per fragment (**out**)
    - At any time in the application
  - Vertex attributes are interpolated by the rasterizer into fragment attributes



# uniform Qualifier

The University of New Mexico

- 
- Variables that are constant for an entire primitive
  - Can be changed in application and sent to shaders
  - Cannot be changed in shader
  - Used to pass information to shader such as the bounding box of a primitive





# in Qualifier

The University of New Mexico

- 
- There are a few built in variables such as `gl_Position` but most have been deprecated
  - User defined (in application program)
    - Use `in` qualifier to get to shader
    - `in float temperature`
    - `in vec3 velocity`



# out Qualifier

The University of New Mexico

- 
- Variables that are passed from vertex shader to fragment shader
  - Automatically interpolated by the rasterizer
  - Old style used the varying qualifier  
`varying vec4 color;`
  - Now use **out** in vertex shader and **in** in the fragment shader  
`out vec4 color;`



# Example: Vertex Shader

The University of New Mexico

---

```
const vec4 red = vec4(1.0, 0.0, 0.0, 1.0);  
in vec4 vPosition;  
out vec3 color_out;  
void main(void)  
{  
    gl_Position = vPosition;  
    color_out = red;  
}
```



# Required Fragment Shader

---

```
in vec3 color_out;  
void main(void)  
{  
    gl_FragColor = color_out;  
}  
  
// in the latest version use form  
// out vec4 fragcolor;  
// fragcolor = color_out;
```



# Passing values

The University of New Mexico

- 
- call by **value-return**
  - Variables are copied in
  - Returned values are copied back
  - Three possibilities
    - **in**
    - **out**
    - **inout** (deprecated)



# Operators and Functions

---

- Standard C functions
  - Trigonometric
  - Arithmetic
  - Normalize, reflect, length
- Overloading of vector and matrix types
  - mat4 a;
  - vec4 b, c, d;
  - $c = b * a$ ; // a column vector stored as a 1d array
  - $d = a * b$ ; // a row vector stored as a 1d array



# Swizzling and Selection

---

- Can refer to array elements by element using [] or selection (.) operator with
  - x, y, z, w
  - r, g, b, a
  - s, t, p, q
  - `a[2]`, `a.b`, `a.z`, `a.p` are the same
- **Swizzling** operator lets us manipulate components

```
vec4 a;  
a.yz = vec2(1.0, 2.0);
```



The University of New Mexico

---

# Programming with OpenGL

## Part 3: Transformations

Ed Angel

Professor of Emeritus of Computer Science  
University of New Mexico





# Objectives

The University of New Mexico

- 
- Learn how to carry out transformations in OpenGL shaders
  - Introduce mat.h and vec.h transformations
    - Model-view
    - Projection



# Transformations and Viewing

---

- In OpenGL, transformation functions are used for changes in coordinate systems
- Projection is also carried out by a projection matrix (transformation)
- Pre 3.0 OpenGL had a set of transformation functions which have been deprecated
- Three choices
  - Application code
  - GLSL functions
  - `vec.h` and `mat.h` (in the textbook)



# Pre 3.1 OpenGL Matrices


---

- In OpenGL matrices were part of the state
- Multiple types
  - Model-View (`GL_MODELVIEW`)
  - Projection (`GL_PROJECTION`)
  - Texture (`GL_TEXTURE`)
  - Color (`GL_COLOR`)
- Single set of functions for manipulation
- Select which to be manipulated by
  - `glMatrixMode(GL_MODELVIEW);`
  - `glMatrixMode(GL_PROJECTION);`

# Current Transformation Matrix (CTM)

---

- Conceptually there is a 4 x 4 homogeneous coordinate matrix, the *current transformation matrix* (CTM) that is part of the state and is applied to all vertices that pass down the pipeline





# Rotation, Translation, Scaling

---

Create an identity matrix:

```
mat4 m = Identity(); // in mat.h
```

Multiply on right by rotation matrix of **theta** in degrees  
where (**vx**, **vy**, **vz**) define axis of rotation

```
mat4 r = Rotate(theta, vx, vy, vz)  
m = m*r;
```

Do same with translation and scaling:

```
mat4 s = Scale( sx, sy, sz)  
mat4 t = Transalate(dx, dy, dz);  
m = m*s*t;
```



# Example

The University of New Mexico

- Rotation about z axis by 30 degrees with a fixed point of (1.0, 2.0, 3.0)

```
mat 4 m = Identity();  
m = Translate(1.0, 2.0, 3.0)*  
    Rotate(30.0, 0.0, 0.0, 1.0)*  
    Translate(-1.0, -2.0, -3.0);
```

- Remember that last matrix specified in the program is the first applied



# Arbitrary Matrices

---

- Can load and multiply by matrices defined in the application program
- Matrices are stored as one dimensional array of 16 elements which are the components of the desired 4 x 4 matrix stored by columns
- OpenGL functions that have matrices as parameters allow the application to send the matrix or its transpose



# Matrix Stacks

The University of New Mexico

- 
- In many situations we want to save transformation matrices for use later
    - Traversing hierarchical data structures (Chapter 8)
    - Avoiding state changes when executing display lists
  - Pre 3.1 OpenGL maintained stacks for each type of matrix
  - Easy to create the same functionality with a simple stack class





# Using Transformations

---

- Example: use idle function to rotate a cube and mouse function to change direction of rotation
- Start with a program that draws a cube in a standard way
  - Centered at origin
  - Sides aligned with axes



# main.c

The University of New Mexico

---

```
void main(int argc, char **argv)
{
    glutInit(&argc, argv);
    glutInitDisplayMode(GLUT_DOUBLE | GLUT_RGB |
        GLUT_DEPTH);
    glutInitWindowSize(500, 500);
    glutCreateWindow("colorcube");
    glutReshapeFunc(myReshape);
    glutDisplayFunc(display);
    glutIdleFunc(spinCube);
    glutMouseFunc(mouse);
    glEnable(GL_DEPTH_TEST);
    glutMainLoop();
}
```



# Idle and Mouse callbacks

The University of New Mexico

```
void spinCube()  
{  
    theta[axis] += 2.0;  
    if( theta[axis] > 360.0 ) theta[axis] -= 360.0;  
    glutPostRedisplay();  
}  
  
void mouse(int btn, int state, int x, int y)  
{  
    if(btn==GLUT_LEFT_BUTTON && state == GLUT_DOWN)  
        axis = 0;  
    if(btn==GLUT_MIDDLE_BUTTON && state == GLUT_DOWN)  
        axis = 1;  
    if(btn==GLUT_RIGHT_BUTTON && state == GLUT_DOWN)  
        axis = 2;  
}
```



# Display callback

The University of New Mexico

---

We can form matrix in application and send to shader and let shader do the rotation or we can send the angle and axis to the shader and let the shader form the transformation matrix and then do the rotation

More efficient than transforming data in application and resending the data

```
void display()  
{  
    glClear(GL_COLOR_BUFFER_BIT | GL_DEPTH_BUFFER_BIT);  
    glUniform(...); //or glUniformMatrix  
    glDrawArrays(...);  
    glutSwapBuffers();  
}
```



# Display callback

The University of New Mexico

```
GLint matrix_loc;
matrix_loc = glGetUniformLocation(program, "rotation");

mat4 ctm;
void display()
{
    glClear(GL_COLOR_BUFFER_BIT | GL_DEPTH_BUFFER_BIT);
    ctm = RotateX(theta[0])*RotateY(theta[1])*RotateZ(theta[2]);
    glUniformMatrix4fv(matrix_loc, 1, GL_TRUE, ctm);
    glDrawArrays(GL_TRIANGLES, 0, N);
    glutSwapBuffers();
}
```



# Vertex Shader


---

```
in vec4 vPosition;  
in vec4 vColor;  
out vec4 color;  
uniform mat4 rotation;  
  
void main()  
{  
    gl_Position = rotation*vPosition;  
    color = vColor;  
}
```



# Transformation Process

- OpenGL had a model-view and a projection matrix in the pipeline which were concatenated together to form the CTM
- We can emulate the process in application
- Vertex shader completes transformation by way of matrix multiplication





# Vertex Shader

The University of New Mexico

---

```
in vec4 vPosition;  
in vec4 vColor;  
out vec4 color;  
uniform mat4 model_view;  
uniform mat4 projection;  
  
void main()  
{  
    gl_Position = projection*model_view*vPosition;  
    color = vColor;  
}
```





# Other Examples

The University of New Mexico

- 
- How to implement?
    - gluLookAt
    - gluOrtho2D
    - glViewport
    - Object animation
    - Camera control (navigation, zoom-in)
  - Many sample codes in the textbook!



The University of New Mexico

---

# Programming with OpenGL

## Part 4: Color and Attributes


Ed Angel

Professor of Emeritus of Computer Science  
University of New Mexico



The University of New Mexico

# OpenGL Primitives







# Polygon Issues

The University of New Mexico

- OpenGL will only display triangles
  - Simple: edges cannot cross
  - Convex: All points on line segment between two points in a polygon are also in the polygon
  - Flat: all vertices are in the same plane
- Application program must tessellate a polygon into triangles (triangulation)
- OpenGL 4.1 contains a tessellator



nonsimple polygon




nonconvex polygon



# Good and Bad Triangles

---

- Long thin triangles render badly



- Equilateral triangles render well
- Maximize minimum angle
- Delaunay triangulation for unstructured points



# Attributes

The University of New Mexico


- 
- Attributes determine the appearance of objects
    - Color (points, lines, polygons)
    - Size and width (points, lines)
    - Stipple pattern (lines, polygons)
    - Polygon mode
      - Display as filled: solid color or stipple pattern
      - Display edges
      - Display vertices
  - Only a few (`glPointSize`) are supported by OpenGL functions



# RGB color

The University of New Mexico

- Each color component is stored separately in the frame buffer
- Usually 8 bits per component in buffer
- Color values can range from 0.0 (none) to 1.0 (all) using floats or over the range from 0 to 255 using unsigned bytels






# Indexed Color

The University of New Mexico

- Colors are indices into tables of RGB values
- Requires less memory
  - indices usually 8 bits
  - not as important now
    - Memory inexpensive
    - Need more colors for shading









# Smooth Color

The University of New Mexico

- Default is *smooth* shading
  - OpenGL interpolates vertex colors across visible polygons
- Alternative is *flat shading*
  - Color of first vertex determines fill color
  - Handle in shader





# Setting Colors

The University of New Mexico

- 
- Colors are ultimately set in the fragment shader but can be determined in either shader or in the application
  - Application color: pass to vertex shader as a uniform variable or as a vertex attribute
  - Vertex shader color: pass to fragment shader as varying variable
  - Fragment color: can alter via shader code



The University of New Mexico

---

# Programming with OpenGL

## Part 5: Complete Programs

Ed Angel

Professor of Emeritus of Computer Science  
University of New Mexico




The University of New Mexico

# A Simple Program (?)

---

Generate a square on a solid background





# It used to be easy

The University of New Mexico

```
#include <GL/glut.h>
void mydisplay(){
    glClear(GL_COLOR_BUFFER_BIT);
    glBegin(GL_QUAD;
        glVertex2f(-0.5, -0.5);
        glVertex2f(-0.5, 0.5);
        glVertex2f(0.5, 0.5);
        glVertex2f(0.5, -0.5);
    glEnd()
}
int main(int argc, char** argv){
    glutCreateWindow("simple");
    glutDisplayFunc(mydisplay);
    glutMainLoop();
}
```



# What happened

The University of New Mexico

- 
- Most OpenGL functions deprecated
  - Makes heavy use of state variable default values that no longer exist
    - Viewing
    - Colors
    - Window parameters
  - New version makes the defaults more explicit
  - However, processing loop is the same



# simple.c

The University of New Mexico

---

```
#include <GL/glut.h>
void mydisplay(){
    glClear(GL_COLOR_BUFFER_BIT);

    // need to fill in this part
    // and add in shaders
}
int main(int argc, char** argv){
    glutCreateWindow("simple");
    glutDisplayFunc(mydisplay);
    glutMainLoop();
}
```



# Immediate Mode Graphics

---

- Geometry specified by vertices
  - Locations in space (2 or 3 dimensional)
  - Points, lines, circles, polygons, curves, surfaces
- Immediate mode
  - Each time a vertex is specified in application, its location is sent to the GPU
  - Old style uses `glVertex`
  - Creates bottleneck between CPU and GPU
  - Removed from OpenGL 3.1






# Retained Mode Graphics

---

- Put all vertex and attribute data in array
- Send array to GPU to be rendered immediately
- Almost OK but problem is we would have to send array over each time we need another render of it
- Better to send array over and store on GPU for multiple renderings



# Display Callback

The University of New Mexico

- Once we get data to GPU, we can initiate the rendering with a simple callback

```
void mydisplay()  
{  
    glClear(GL_COLOR_BUFFER_BIT);  
    glDrawArrays(GL_TRIANGLES, 0, 3);  
    glFlush();  
}
```

- Arrays are buffer objects that contain vertex arrays



# Vertex Arrays

The University of New Mexico

- Vertices can have many attributes
  - Position
  - Color
  - Texture Coordinates
  - Application data
- A vertex array holds these data
- Using types in `vec.h`

```
point2 vertices[3] = {point2(0.0, 0.0),  
                      point2(0.0, 1.0), point2(1.0, 1.0)};
```



# Vertex Array Object

The University of New Mexico

- Bundles all vertex data (positions, colors, ...)
- Get name for buffer then bind

```
Glunit abuffer;  
glGenVertexArrays(1, &abuffer);  
glBindVertexArray(abuffer);
```

- At this point we have a current vertex array but no contents
- Use of glBindVertexArray lets us switch between VBOs



# Buffer Object

The University of New Mexico

- Buffer objects allow us to transfer large amounts of data to the GPU
- Need to create, bind and identify data

```
Gluint buffer;  
glGenBuffers(1, &buffer);  
glBindBuffer(GL_ARRAY_BUFFER, buffer);  
glBufferData(GL_ARRAY_BUFFER,  
             sizeof(points), points);
```

- Data in current vertex array is sent to GPU



# More Information

The University of New Mexico

- 
- VAOs (Vertex Array Objects)
  - glBindVertexArrays vs glBindBuffer



# Initialization

The University of New Mexico

- 
- Vertex array objects and buffer objects can be set up on **init()**
  - Also set clear color and other OpenGL parameters
  - Also set up shaders as part of initialization
    - Read
    - Compile
    - Link



The University of New Mexico

# Linking Shaders with Application

---

- Read shaders
- Compile shaders
- Create a program object
- Link everything together
- Link variables in application with variables in shaders
  - Vertex attributes
  - Uniform variables





# Program Object

The University of New Mexico

- Container for shaders
  - Can contain multiple shaders
  - Other GLSL functions

```
GLuint myProgObj;  
myProgObj = glCreateProgram();  
/* define shader objects here */  
glUseProgram(myProgObj);  
glLinkProgram(myProgObj);
```



# Reading a Shader

The University of New Mexico

- 
- Shaders are added to the program object and compiled
  - Usual method of passing a shader is as a null-terminated string using the function **glShaderSource**
  - If the shader is in a file, we can write a reader to convert the file to a string
  - Sample code in the textbook



# Adding a Vertex Shader

---

```
GLuint vShader;  
GLuint myVertexObj;  
GLchar vShaderfile[] = "my_vertex_shader";  
GLchar* vSource =  
    readShaderSource(vShaderFile);  
myVertexObj =  
    glCreateShader(GL_VERTEX_SHADER);  
glShaderSource(myVertexObj, 1,  
    &vSource, NULL);  
glCompileShader(myVertexObj);  
glAttachObject(myProgObj, myVertexObj);
```



# Vertex Attributes

The University of New Mexico

- 
- Vertex attributes are named in the shaders
  - Linker forms a table
  - Application can get index from table and tie it to an application variable
  - Similar process for uniform variables



# Vertex Attribute Example

---

```
#define BUFFER_OFFSET( offset )  
    ((GLvoid*) (offset))
```

```
GLuint loc =  
    glGetAttribLocation( program, "vPosition" );  
glEnableVertexAttribArray( loc );  
glVertexAttribPointer( loc, 2, GL_FLOAT,  
    GL_FALSE, 0, BUFFER_OFFSET(0) );
```



# Uniform Variable Example

---

```
GLint angleParam;  
angleParam = glGetUniformLocation(myProgObj,  
    "angle");  
/* angle defined in shader */  
  
/* my_angle set in application */  
GLfloat my_angle;  
my_angle = 5.0 /* or some other value */  
  
glUniform1f(angleParam, my_angle);
```



# Adding Color

The University of New Mexico

- 
- If we set a color in the application, we can send it to the shaders as a vertex attribute or as a uniform variable depending on how often it changes
  - Let's associate a color with each vertex
  - Set up an array of same size as positions
  - Send to GPU as a vertex buffer object



# Setting Colors

The University of New Mexico

---

```
typedef vec3 color3;  
color3 base_colors[4] = {color3(1.0, 0.0, 0.0), ....  
color3 colors[NumVertices];  
vec3 points[NumVertices];
```

```
//in loop setting positions
```

```
colors[i] = basecolors[color_index]  
position[i] = .....
```





# Setting Up Buffer Object

---

```
//need larger buffer
```

```
glBufferData(GL_ARRAY_BUFFER, sizeof(points) +  
    sizeof(colors), NULL, GL_STATIC_DRAW);
```

```
//load data separately
```

```
glBufferSubData(GL_ARRAY_BUFFER, 0,  
    sizeof(points), points);  
glBufferSubData(GL_ARRAY_BUFFER, sizeof(points),  
    sizeof(colors), colors);
```



# Second Vertex Array

The University of New Mexico

---

// vPosition and vColor identifiers in vertex shader

```
loc = glGetAttribLocation(program, "vPosition");  
glEnableVertexAttribArray(loc);  
glVertexAttribPointer(loc, 3, GL_FLOAT, GL_FALSE, 0,  
    BUFFER_OFFSET(0));
```

```
loc2 = glGetAttribLocation(program, "vColor");  
glEnableVertexAttribArray(loc2);  
glVertexAttribPointer(loc2, 3, GL_FLOAT, GL_FALSE, 0,  
    BUFFER_OFFSET(sizeofpoints));
```



The University of New Mexico

---

# Shader Applications

Ed Angel

Professor Emeritus of Computer Science

University of New Mexico



# Vertex Shader Applications

---

- Moving vertices
  - Wave motion
  - Particle system
  - Morphing
  - Fractals
- Lighting
  - More realistic models
  - Cartoon shaders



# Wave Motion Vertex Shader

---

```
in vec4 vPosition;
uniform float xs, zs; // frequencies
uniform float h; // height scale
void main()
{
    vec4 t = vPosition;
    t.y = vPosition.y
        + h*sin(time + xs*vPosition.x)
        + h*sin(time + zs*vPosition.z);
    gl_Position = t;
}
```



# Particle System

The University of New Mexico

```
in vec3 vPosition;
uniform mat4 ModelViewProjectionMatrix;
uniform vec3 init_vel;
uniform float g, m, t;

void main()
{
    vec3 object_pos;
    object_pos.x = vPosition.x + vel.x*t;
    object_pos.y = vPosition.y + vel.y*t
                + g/(2.0*m)*t*t;
    object_pos.z = vPosition.z + vel.z*t;
    gl_Position = ModelViewProjectionMatrix
                *vec4(object_pos,1);
}
```



# Pass Through Fragment Shader

---

```
/* pass-through fragment shader */
```

```
in vec4 color;  
void main(void)  
{  
    gl_FragColor = color;  
}
```



The University of New Mexico

# Vertex vs Fragment Lighting

---



per vertex lighting



per fragment lighting






The University of New Mexico

# Fragment Shader Applications

---


## Texture mapping



smooth shading



environment  
mapping



bump mapping



# Summary

The University of New Mexico

- 
- Programmable graphics pipeline
  - Vertex processing → vertex shader
  - Fragment process → fragment shader
  - Binding shaders to application program
  - Transferring values to shaders
  - Many sample codes in the textbook!
    - vec.h and mat.h in the book web page



# Summary

The University of New Mexico

- 
- Lots of advanced features of shaders
  - Geometry/Tessellation/Compute shaders
  - Multiple passes of shaders
  - Massive parallelism → high performance
  - Used for a variety of applications
  - GPGPU (General-Purpose GPU)
    - CUDA, OpenCL



The University of New Mexico

---

# Supplementary Slides



# OpenGL Libraries

The University of New Mexico

- 
- OpenGL core library
    - OpenGL32 on Windows
    - GL on most unix/linux systems (libGL.a)
  - OpenGL Utility Library (GLU)
    - Provides functionality in OpenGL core but avoids having to rewrite code
    - Will only work with legacy code
  - Links with window system
    - GLX for X window systems
    - WGL for Windows
    - AGL for Macintosh



- 
- OpenGL Utility Toolkit (GLUT)
    - Provides functionality common to all window systems
      - Open a window
      - Get input from mouse and keyboard
      - Menus
      - Event-driven
    - Code is portable but GLUT lacks the functionality of a good toolkit for a specific platform
      - No slide bars



# freeglut

The University of New Mexico

- 
- GLUT was created long ago and has been unchanged
    - Amazing that it works with OpenGL 3.1
    - Some functionality can't work since it requires deprecated functions
  - freeglut updates GLUT
    - Added capabilities
    - Context checking



- 
- OpenGL Extension Wrangler Library
  - Makes it easy to access OpenGL extensions available on a particular system
  - Avoids having to have specific entry points in Windows code
  - Application needs only to include `glew.h` and run a `glewInit()`






The University of New Mexico

# Software Organization

---





# OpenGL #defines

---

- Most constants are defined in the include files `gl.h`, `glu.h` and `glut.h`
  - Note `#include <GL/glut.h>` should automatically include the others
  - Examples
    - `glEnable(GL_DEPTH_TEST)`
    - `glClear(GL_COLOR_BUFFER_BIT)`
- include files also define OpenGL data types: `GLfloat`, `GLdouble`, ....



# OpenGL Functions

The University of New Mexico

- 
- Primitives
    - Points
    - Line Segments
    - Triangles
  - Attributes
  - Transformations
    - Viewing
    - Modeling
  - Control (GLUT)
  - Input (GLUT)
  - Query



The University of New Mexico

# OpenGL function format

---

function name      dimensions

`glUniform3f(x, y, z)`

belongs to GL library      `x, y, z` are floats

`glUniform3fv(p)`

`p` is a pointer to an array



# Lack of Object Orientation

---

- OpenGL is not object oriented so that there are multiple functions for a given logical function
  - glUniform3f
  - glUniform2i
  - glUniform3dv
- Easy to create overloaded functions in C++ but issue is efficiency



# OpenGL State

The University of New Mexico

- 
- OpenGL is a state machine
  - OpenGL functions are of two types
    - Primitive generating
      - Can cause output if primitive is visible
      - Vertices are processed and appearance of primitive are controlled by the state
    - State changing
      - Transformation functions
      - Attribute functions
      - Under 3.1 most state variables are defined by the application and sent to the shaders



# main.c

The University of New Mexico

```
#include <GL/glew.h>
```

```
#include <GL/glut.h>
```

includes **gl.h**

```
int main(int argc, char** argv)
```

```
{
```

```
    glutInit(&argc, argv);
```

```
    glutInitDisplayMode(GLUT_SINGLE | GLUT_RGB);
```

```
    glutInitWindowSize(500, 500);
```

```
    glutInitWindowPosition(0, 0);
```

```
    glutCreateWindow("simple");
```

```
    glutDisplayFunc(mydisplay);
```

```
    glewInit();
```

```
    init();
```

```
    glutMainLoop();
```

```
}
```

enter event loop



# GLUT functions

The University of New Mexico

- 
- **glutInit** allows application to get command line arguments and initializes system
  - **gluInitDisplayMode** requests properties for the window (the *rendering context*)
    - RGB color
    - Single buffering
    - Properties logically ORed together
  - **glutWindowSize** in pixels
  - **glutWindowPosition** from top-left corner of display
  - **glutCreateWindow** create window with title “simple”
  - **glutDisplayFunc** display callback
  - **glutMainLoop** enter infinite event loop





# Double Buffering

The University of New Mexico

- 
- Updating the value of a uniform variable opens the door to animating an application
    - Execute glUniform in display callback
    - Force a redraw through glutPostRedisplay()
  - Need to prevent a partially redrawn frame buffer from being displayed
  - Draw into back buffer
  - Display front buffer
  - Swap buffers after updating finished



# Adding Double Buffering

---

- Request a double buffer
  - `glutInitDisplayMode(GLUT_DOUBLE)`
- Swap buffers

```
void mydisplay()
{
    glClear(.....);
    glDrawArrays();
    glutSwapBuffers();
}
```



# Event Loop

The University of New Mexico

- 
- Note that the program specifies a *display callback* function named **mydisplay**
    - Every glut program must have a display callback
    - The display callback is executed whenever OpenGL decides the display must be refreshed, for example when the window is opened
    - The **main** function ends with the program entering an event loop



# Idle Callback

The University of New Mexico

- Idle callback specifies function to be executed when no other actions pending
  - `glutIdleFunc(myIdle);`

```
void myIdle()
{
    // recompute display
    glutPostRedisplay();
}
```



# Program Structure

---

- Most OpenGL programs have a similar structure that consists of the following functions

- **main( )**:

- specifies the callback functions
- opens one or more windows with the required properties
- enters event loop (last executable statement)

- **init( )**: sets the state variables

- Viewing
- Attributes

- **initShader( )**: read, compile and link shaders

- **callbacks**

- Display function
- Input and window functions



# simple.c revisited

The University of New Mexico

- 
- **main( )** function similar to last lecture
    - Mostly GLUT functions
  - **init()** will allow more flexible colors
  - **initShader()** will hide details of setting up shaders for now
  - Key issue is that we must form a data array to send to GPU and then render it



# Shader Reader

The University of New Mexico

```
#include <stdio.h>
```

```
static char*
```

```
readShaderSource(const char* shaderFile)
```

```
{
```

```
    FILE* fp = fopen(shaderFile, "r");
```

```
    if ( fp == NULL ) { return NULL; }
```

```
    fseek(fp, 0L, SEEK_END);
```

```
    long size = ftell(fp);
```



# Shader Reader (cont)

The University of New Mexico

```
fseek(fp, 0L, SEEK_SET);  
char* buf = new char[size + 1];  
fread(buf, 1, size, fp);
```

```
buf[size] = '\0';  
fclose(fp);
```

```
return buf;  
}
```





# Attribute and Varying Qualifiers

---

- Starting with GLSL 1.5 attribute and varying qualifiers have been replaced by in and out qualifiers
- No changes needed in application
- Vertex shader example:

```
#version 1.4
attribute vec3 vPosition;
varying vec3 color;
```

```
#version 1.5
in vec3 vPosition;
out vec3 color;
```



# Coordinate Systems

The University of New Mexico


- 
- The units in **points** are determined by the application and are called *world*, *object*, *model* or *problem coordinates*
  - Viewing specifications usually are also in object coordinates
  - Eventually pixels will be produced in *device*, *window coordinates*
  - OpenGL also uses some internal representations that usually are not visible to the application but are important in the shaders



# OpenGL Camera

The University of New Mexico

- OpenGL places a camera at the origin in object space pointing in the negative  $z$  direction
- The default viewing volume is a box centered at the origin with sides of length 2







# Orthographic Viewing

The University of New Mexico

In the default orthographic view, points are projected forward along the  $z$  axis onto the plane  $z=0$






# Viewports

The University of New Mexico

- Do not have to use the entire window for the image: `glViewport(x, y, w, h)`
- Values in pixels (device coordinates)





# Polygon Testing


The University of New Mexico

- 
- Conceptually simple to test for simplicity and convexity
  - Time consuming
  - Earlier versions assumed both and left testing to the application
  - Present version only renders triangles
  - Need algorithm to triangulate an arbitrary polygon



# Triangularization

- Convex polygon




- Start with abc, remove b, then acd, ....



The University of New Mexico

# Non-convex (concave)

---








# Recursive Division

The University of New Mexico

- Find leftmost vertex and split





# CTM operations

The University of New Mexico

- The CTM can be altered either by loading a new CTM or by postmultiplication

Load an identity matrix:  $\mathbf{C} \leftarrow \mathbf{I}$

Load an arbitrary matrix:  $\mathbf{C} \leftarrow \mathbf{M}$

Load a translation matrix:  $\mathbf{C} \leftarrow \mathbf{T}$

Load a rotation matrix:  $\mathbf{C} \leftarrow \mathbf{R}$

Load a scaling matrix:  $\mathbf{C} \leftarrow \mathbf{S}$

Postmultiply by an arbitrary matrix:  $\mathbf{C} \leftarrow \mathbf{C}\mathbf{M}$

Postmultiply by a translation matrix:  $\mathbf{C} \leftarrow \mathbf{C}\mathbf{T}$

Postmultiply by a rotation matrix:  $\mathbf{C} \leftarrow \mathbf{C}\mathbf{R}$

Postmultiply by a scaling matrix:  $\mathbf{C} \leftarrow \mathbf{C}\mathbf{S}$



# Rotation about a Fixed Point

---

Start with identity matrix:  $C \leftarrow I$

Move fixed point to origin:  $C \leftarrow CT$

Rotate:  $C \leftarrow CR$

Move fixed point back:  $C \leftarrow CT^{-1}$

Result:  $C = TRT^{-1}$  which is **backwards**.

This result is a consequence of doing postmultiplications.  
Let's try again.



# Reversing the Order

The University of New Mexico

---

We want  $\mathbf{C} = \mathbf{T}^{-1} \mathbf{R} \mathbf{T}$   
so we must do the operations in the following order

$$\mathbf{C} \leftarrow \mathbf{I}$$

$$\mathbf{C} \leftarrow \mathbf{C} \mathbf{T}^{-1}$$

$$\mathbf{C} \leftarrow \mathbf{C} \mathbf{R}$$

$$\mathbf{C} \leftarrow \mathbf{C} \mathbf{T}$$

Each operation corresponds to one function call in the program.

Note that the last operation specified is the first executed in the program